
Q-Learning vs. DQN

Comparing learning algorithms on Atari game Space Invaders

— By Divyam Garg & John DiNofrio —

Reinforcement Learning

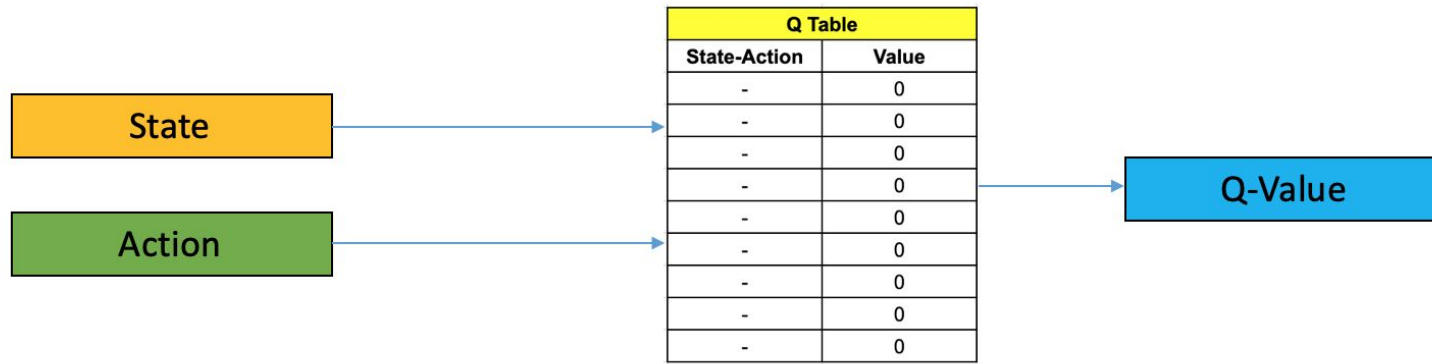
- Reinforcement learning is one of the three most basic machine learning techniques
- Unlike supervised and unsupervised learning
 - It does not require labeled data
 - It does not need to directly correct every action
- By setting positive and negative rewards, the program can teach itself how to play games or perform actions




Space Invaders

- We used it to teach the computer how to play the Atari game **Space Invaders**
- We used both Q-Learning and DQN to demonstrate this and compare how the two methods worked





Q Learning



Deep Q Learning

Q-Learning

- Q-Learning is quite simple to implement
- The entire algorithm depends on an array called a **Q-Table**
- This Q-Table is $n \times m$ dimensions
 - Where n is the number of actions possible
 - Where m is the number of states possible
- For every current state and action that led it there, the value in the Q-Table will change
 - The value will go up if the reward was positive for that action
 - The value will go down if the reward was negative
 - The value could go up or down if there is no reward - this depends on the parameters you set

Q-table initialised at zero

	UP	DOWN	LEFT	RIGHT
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0
5	0	0	0	0
6	0	0	0	0
7	0	0	0	0
8	0	0	0	0

After few episodes

	UP	DOWN	LEFT	RIGHT
0	0	0	0	0
1	0	0	0	0
2	0	2.25	2.25	0
3	0	0	5	0
4	0	0	0	0
5	0	0	0	0
6	0	5	0	0
7	0	0	2.25	0
8	0	0	0	0

Eventually

	UP	DOWN	LEFT	RIGHT
0	0	0	0.45	0
1	0	1.01	0	0
2	0	2.25	2.25	0
3	0	0	5	0
4	0	0	0	0
5	0	0	0	0
6	0	5	0	0
7	0	0	2.25	0
8	0	0	0	0

Q-Learning


- For both Q-Learning and DQN, there are certain parameters you can set
 - **Learning rate** - how much new data overwrites old data
 - 0 = no learning 1 = only new info is retained
 - **Discount rate** - short term or long term goals
 - 0 = short-sighted 1 = far-sighted
 - **Exploration rate** - how often the program tries random, new actions
 - 0 = never 1 = always new and random
 - **Number of episodes** - the amount of attempts it gets (how many games)
 - **Number of steps** - the number of actions it can take per episode

Q-Learning on Space Invaders


- **Learning rate** - 0.3
- **Discount rate** - 0.7
- **Exploration rate** - Started at 1 and exponentially went down
- **Number of episodes** - the amount of attempts it gets (how many games)
- **Number of steps** - Max 50000
- **Action space** - 6 (move left, move right, shoot, move left/shoot, move right/shoot, do nothing)
- **State space**- 160 (width of game)

Q-Learning Training

Single Game Score vs Step Number for Q-Learning Training




Total Reward vs Step Number for Q-Learning Training



Q-Learning Test

200 Episodes



Q-Learning Test

200 Episodes

Score 50.0

Score 160.0

Score 210.0

Score 155.0

Score 180.0

Score 155.0

Score 80.0

Score 210.0

Score 210.0

Score 180.0


Average Score: 159.0




Q-Learning Training

200 Episodes

Single Game Score vs Step Number for Q-Learning Training

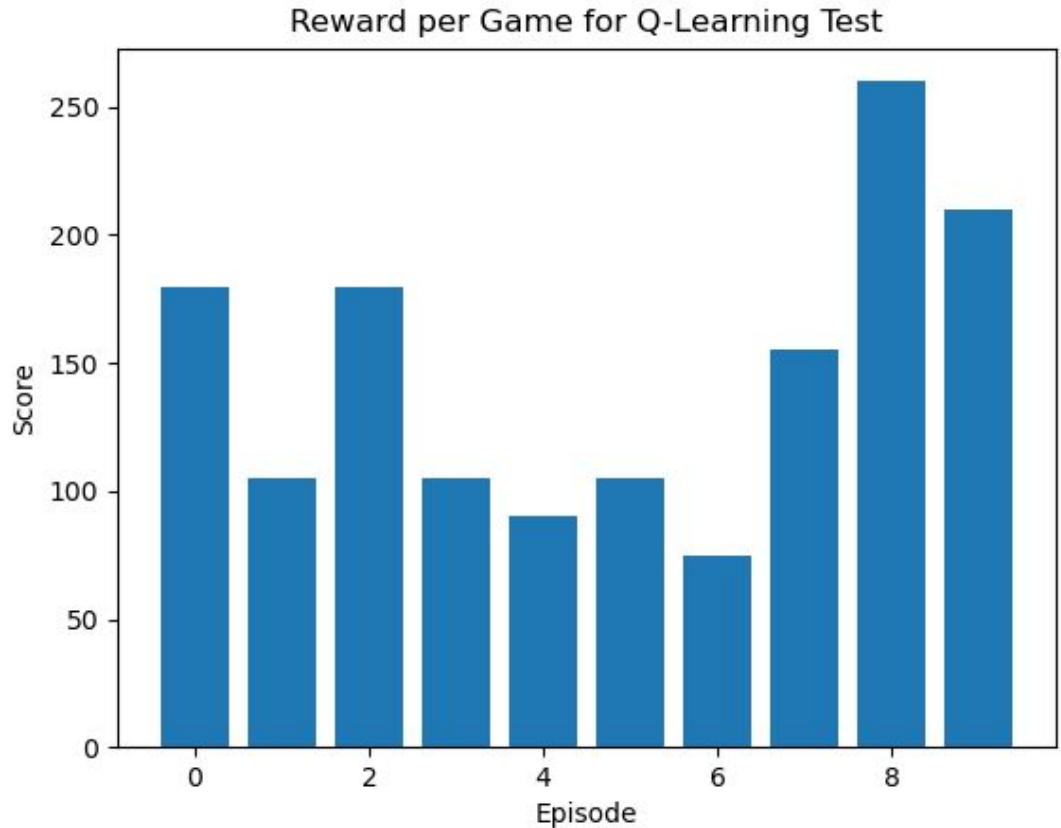


Total Reward vs Step Number for Q-Learning Training



Q-Learning Test

200 Episodes



Q-Learning Test

500 Episodes

Score 180.0

Score 105.0

Score 180.0

Score 105.0

Score 90.0

Score 105.0

Score 75.0

Score 155.0

Score 260.0

Score 210.0

Average Score: 146.5




Q-Learning Results

- Unfortunately, the Q-Learning algorithm did not perform very well. It learned the basics and could get a few points in before it died, but it didn't *play* the game
- The reason for this is the state space
 - The state space for Q-Learning was the position of the spaceship
 - It had no idea where the bullets or ships were
- What it did learn
 - With the little information it had, it knew that it wouldn't get hurt while hiding under the shield
 - It also knew that if it kept shooting with half its body protected away it could get the most points
 - Q-Learning doesn't have the capacity to process that much information and fine tune its training

Overcoming Shortcomings with Deep - Q Learning

- The problem with games is that the agent must take in images as perception; therefore, we need to use a convolutional neural network to add precision.
- The state space for DQN is a lot bigger, and it includes every pixel on the screen
 - This means that DQN knows where every bullet and enemy is as well as how much shield is left



Implementation of Deep Q-learning Algorithm

- Initialize the environment of the Space-Invaders-Atari 2600 (8 actions possible)
- Preprocessing the frame:
 - Removal of unnecessary pixel (210,160) to (110,84)
 - Normalise for better distribution and resizing the image
- Getting stack frames
- Setting up training and hyper-parameters (state size : - 110,84,4)
 - Learning rate: - 0.0025 & decay rate = 0.0001
 - Total episodes = 150
 - Batch size = 64
 - Explore probability at start = 1.0
 - Discount = 0.9
 - Memory size = 1000000

Making it work

- Memory sample: - Storing previous experiences
- Replay: - Avoid reinforcing the same experience
- Creating the hidden layers
 - 3 Conv2d layers activation = 'elu'
 - 1 flatten
 - 2 Dense layers activation = 'softmax'
- Predicting the future Q (action = $\text{argmax}Q(s,a)$)



Deep Q-learning training

$$\Delta w = \alpha [(\underbrace{R + \gamma \max_a \hat{Q}(s', a, w)}_{\text{Maximum possible Qvalue for the next_state (= Q_target)}}) - \underbrace{\hat{Q}(s, a, w)}_{\text{Current predicted Q-val}}] \nabla_w \hat{Q}(s, a, w)$$

Change in
weights

learning
rate


Maximum possible Qvalue for the
next_state (= Q_target)

Current predicted
Q-val

TD Error

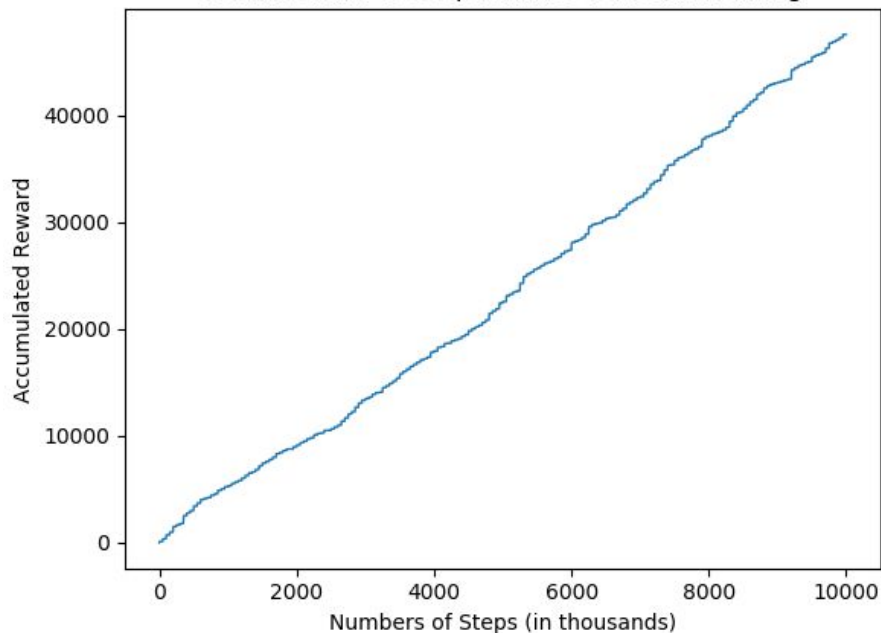
Gradient of our current
predicted Q-value

DQN Training - 50 Episodes




DQN Training (200 Episodes) vs Q-Learning (500 Episodes)

Total Reward vs Step Number for DQN Training



Total Reward vs Step Number for Q-Learning Training



DQN Results - Training 200 Episodes/Test 10 Games

Score 480.0

Score 520.0

Score 260.0

Score 405.0

Score 190.0

Score 855.0

Score 630.0


Score 345.0


Score 305.0

Score 225.0

Average Score: 394.

Score per Episode DQN Training





DQN Results

- DQN performed a lot better than Q-Learning did, but it took a lot more time to train
 - Top Score - 855 (compared to Q-Learning's 260)
 - Avg Score - 394 (compared to Q-Learning's 146)
- With a very large state space, DQN could find the best move for every scenario
- This also meant that the training time was a lot longer
 - Training just 50 episodes took almost 3 hours
 - Training 50 episodes with Q-Learning only took about 10 minutes
- The DQN algorithm could have performed even better if we gave it more time to train
 - Q-Learning had a lot fewer steps (by a factor of 100) but a lot more episodes (300 more episodes), and it still did not do as well
- When it comes down to it, DQN actually played the game
 - It knew where to hide, when to shoot, and the pink think at the top is bonus points
 - Q-Learning had no clue about any of that except how to hide

References

<https://www.freecodecamp.org/news/an-introduction-to-reinforcement-learning-4339519de419/>

https://simoniniithomas.github.io/Deep_reinforcement_learning_Course/

<https://yilundu.github.io/2016/12/24/Deep-Q-Learning-on-Space-Invaders.html>

https://www.youtube.com/watch?v=PnHCvfgC_ZA&list=PL7-jPKtc4r78-wCZcQn5lqyuWhBZ8fOxT&index=4

https://www.youtube.com/watch?v=0g4j2k_Ggc4&list=PL7-jPKtc4r78-wCZcQn5lqyuWhBZ8fOxT&index=5

<https://www.youtube.com/watch?v=gCJyVX98KJ4>